

# Diagnostics

Subrata Paul

6/4/2020

# What we need to check?

- Nonlinearity of regression function
- Nonconstancy of error variance
- Nonindependence of error terms
- Nonnormality of error terms
- Presence of outliers and influential observations

# Residual Plots


# Importance

- Diagnostic plots of the response variable are not very useful
- Residual plots usually (indirectly) carries out diagnostics of the response variable
- Residuals gives us information about model structure and the nature of unobserved randomness.


# Different plots of residuals

- Plots of residuals against predictor variables
- Plot of residuals against fitted values
- Plot of absolute or squared residuals against fitted values or predictor variables
- Plot of residuals against time or other sequence
- Plots of residuals against omitted predictor variables
- Box plot and density plot of residuals
- Q-Q plot of residuals


# Reminder




# Simulated scenarios



(a)



(b)




(c)



(d)

# Example (Non-constant error variance)





# Test for non-constant variance

- If the data is separable in some way then a F test works
- Brown\_Forsythe test (a t-test)
- Breusch-Pagan test (also known as Cook-Weisberg score test)
  - It is a chi-squared test

# Residual Plots

- It can be used to check
  - Nonlinear structure
  - Non-constant error variance
  - Non-zero error mean
- Though it is not used directly to detect unusual observations but gives us an overview of the observations.
- Though it indicates potential structural problem it does not clearly suggest transformation because it fails to distinguish between monotonic and non-monotonic transformation.
- Does not show the nature of the marginal effect of a predictor variable, given the other predictor variables in the model.


# Interesting Extreme Example

```
library(faraway)
set.seed(123)
male = 600
female = 400

bp = c(rnorm(male, 85, 2), rnorm(female, 75, 1))
bmi = c(rnorm(male, 24, 4), rnorm(female, 28, 2))
chol = c(rnorm(male, 200, 10), rnorm(female, 150, 15))
gender = c(rep(1, male), rep(0, female))
dat = data.frame(bp = bp, bmi = bmi, chol = chol, gender = gender)
lmod = lm(bmi ~ bp + chol, data = dat)
summary(lmod)

##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 48.4499279  2.1428433 22.6101 < 2.2e-16
## bp          -0.2169564  0.0388197 -5.5888  2.95e-08
## chol        -0.0286526  0.0074002 -3.8719  0.000115
##
## n = 1000, p = 3, Residual SE = 3.38723, R-Squared = 0.23
```

# Interesting Extreme Example (Residual Plot)



# What does the residual plot suggests?

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9.7006468   7.8100284   1.2421  0.21470
## bp          0.1604810   0.0823442   1.9489  0.05177
## chol        0.0038192   0.0168418   0.2268  0.82068
##
## n = 600, p = 3, Residual SE = 3.89664, R-Squared = 0.01

##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 12.0170106   7.8054019   1.5396  0.12446
## bp          0.2005841   0.1027723   1.9517  0.05167
## chol        0.0071092   0.0068681   1.0351  0.30125
##
## n = 400, p = 3, Residual SE = 2.11171, R-Squared = 0.01
```

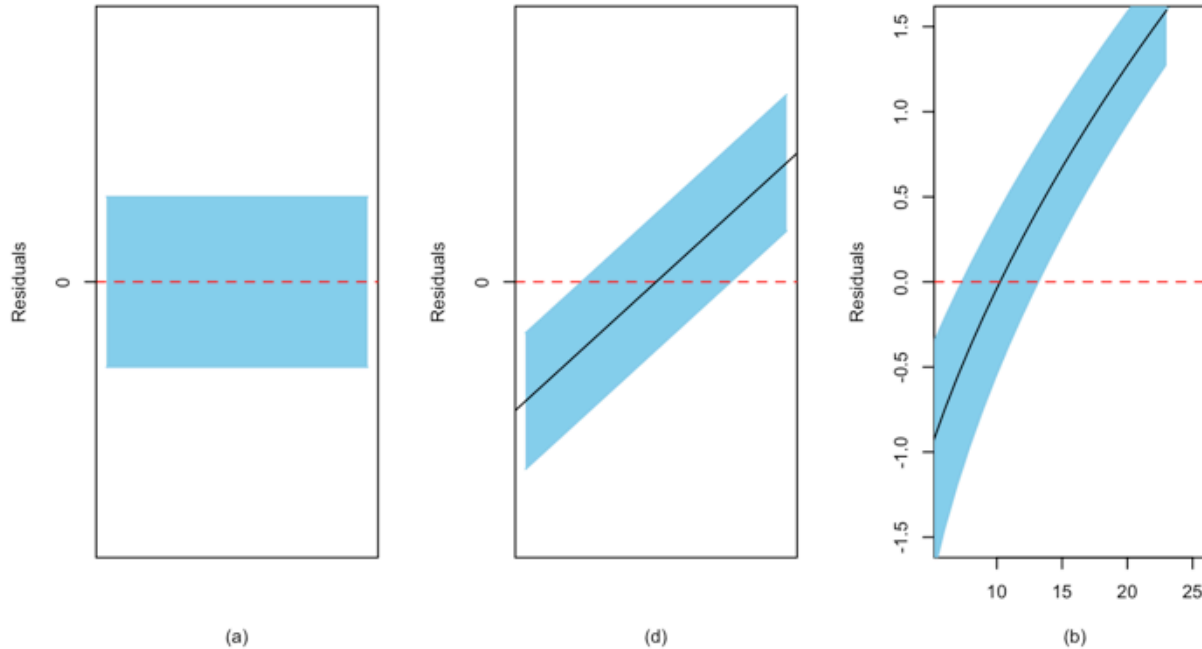
# Test for non-constant variance

- If the data is separable in some way then a F test works
- Brown\_Forsythe test (a t-test)
- Breusch-Pagan test (also known as Cook-Weisberg score test)
  - It is a chi-squared test

# Added variable plot

- Gives marginal relationship given the other predictors are in the model
- Can be used to investigate if a new variable is worthy to include in the model
- Provide clearer picture compared to residual vs predictor plot


# Simulated Scenarios




1. No additional information from the predictor
2. Helpful addition of the predictor
3. Inclusion of the predictor with some transformation.



# AV plot example



**Added-Variable Plot: Income**



# Component plus residual (CR) plot

The **component plus residual (cr) plot** (a.k.a, **partial residual plot**) is a competitor to the added variable plot.

The **cr plot** shows  $\hat{\beta}_i x_i + \hat{e}$  versus  $x_i$ .

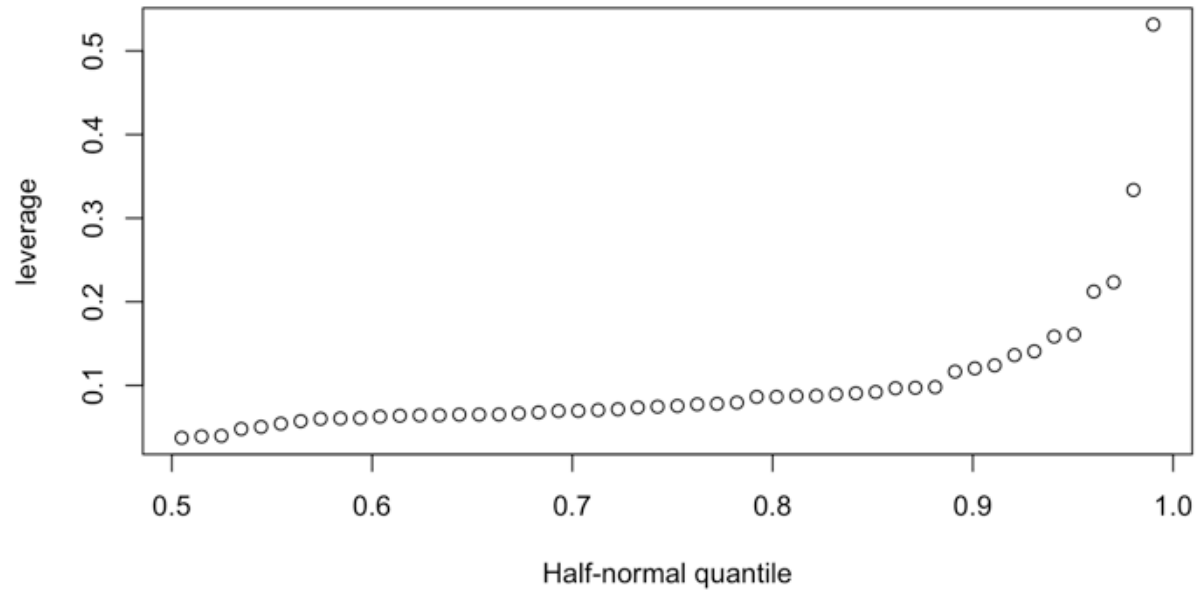
- **cr** plots are useful for checking nonlinear relationships in the variable being considered for inclusion in the model.
- They can also suggest potential transformation of the data so that the relationship is linear.
- If the scatter plot does not appear to be linear, then there is a nonlinear relationship between the regressor and the response (after accounting for the other regressors).
- The slope of the line fit to the cr plot is  $\hat{\beta}_i$ .

Unusual observations

# Leverage

- A **leverage** point is an observation that is unusual in the predictor space.
- $h_{ii}$  is called the leverage value of the  $i$  th observation.
- $h_{ii}$  is a measure of distance between the  $X$  values for the  $i$ th observation from the mean of the  $X$  values of all  $n$  observations.
- $h_{ii}$  measures the role of the  $X$  values in determining how important  $y_i$  is in affecting  $\hat{y}_i$
- A half-normal plot of the leverage values can be used to identify observations with unusually high leverage.
- A leverage value  $h_i$  is usually considered large if it is more than twice as large as the mean leverage value.

# Example



# Outlier

An outlier is a point that does not fit the current model.


- An outlier is context specific! An outlier for one model may not be an outlier for a different model.

externally **studentized** residual

$$t_i = \frac{y_i - \hat{y}_{(i)}}{\hat{\sigma}_{(i)} \sqrt{1 + x_i^T (X_{(i)}^T X_{(i)})^{-1} x_i}} \sim T_{n-p-1}.$$

- Bonferroni correction is needed

# Example



# Influential observations

- An influential observation is one whose removal from the dataset would cause a large change in the fitted model.
- An influential observation is usually a leverage point, an outlier, or both.


The Cook's distance is a popular inferential tool because it reduces influence information to a single value for each observation.

The Cook's distance for the  $i$ th observation is


$$D_i = \frac{(\hat{\mathbf{y}} - \hat{\mathbf{y}}_{(i)})^T (\hat{\mathbf{y}} - \hat{\mathbf{y}}_{(i)})}{p \hat{\sigma}^2} = \frac{1}{p} r_i^2 \frac{h_i}{1 - h_i}$$



# Example



# Influence plot



```
##           StudRes      Hat      CookD
## Chile      -2.3134295 0.03729796 0.03781324
## Japan       1.6032158 0.22330989 0.14281625
## United States -0.3546151 0.33368800 0.01284481
## Zambia      2.8535583 0.06433163 0.09663275
```

# Correct or Delete the Observation(s)

- If they're data entry errors, correct the problem. If they can't be fixed, remove them (they're wrong, so they don't tell us anything useful).
- Remove them if they're not part of the population of interest (you are studying dogs, but this observation is a cat).
- Remove them because they break the model.
- This is a bad idea.
- Make sure to indicate that you removed them from the data set and explain why.
- *THIS IS A BAD IDEA.*

# Fit a Different Model

- An outlier/influential point for one model may not be for another.
- Examine the physical context—why did it happen?
- An outlier/influential point may be interesting in itself.
  - An outlier in a statistical analysis of credit card transactions may indicate fraud!
- This may suggest a better model.
- Use robust regression, which is not as affected by outliers/influential observations.
- Never automatically remove outliers/influential points!
- They provide important information that may otherwise be missed.
- Fit the model with and without the influential observation(s).
- Do your results substantively change?

# Checking Error

Summary of methods for checking error assumptions

- Mean-zero error assumption:
  - Plot of residuals versus fitted values
- Constant error variance assumption:
  - Plot of residuals versus fitted values
  - Plot of  $\sqrt{|\hat{\epsilon}|}$  versus fitted values.
- Normal error assumption:
  - q-q of residuals
  - Shapiro-wilk test
- Autocorrelated errors:
  - Plot of residuals versus time
  - Plot of successive pairs of residuals
  - Durbin-Watson test

# Weighted Least Squares

# Model

The generalized multiple regression model:

$$Y_i = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_{p-1} X_{i,p-1} + \epsilon_i$$

where,

- $\beta_0, \beta_1, \dots, \beta_{p-1}$  are parameters
- $X_{i1}, \dots, X_{i,p-1}$  are known constants
- $\epsilon_i$  are independent  $N(0, \sigma_i^2)$

# Variance-covariance matrix

$$Var(\epsilon) = \begin{bmatrix} \sigma_1^2 & 0 & \dots & 0 \\ 0 & \sigma_2^2 & \dots & 0 \\ \vdots & \vdots & & \vdots \\ 0 & 0 & \dots & \sigma_n^2 \end{bmatrix}$$

- OLS assumes equal variance:  $\sigma_1^2 = \dots = \sigma_n^2 = \sigma^2$
- Using OLS we would get unbiased estimation of the parameters
- The OLS estimates no longer have minimum variance
- We must account for unequal variance in the estimation process
- Consider three cases:
  - Error variances are known (unrealistic)
  - Error variances are known up to proportionality constant
  - Error variances are known (realistic)



# Error variances are known

Likelihood

$$L(\beta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma_i^2}} \exp \left[ -\frac{1}{2\sigma_i^2} (Y_i - \beta_0 - \beta_1 X_{i1} - \dots - \beta_{p-1} X_{i,p-1})^2 \right]$$

Define

$$w_i = \frac{1}{\sigma_i^2}$$

$$L(\beta) = \left[ \prod_{i=1}^n \frac{\sqrt{w_i}}{\sqrt{2\pi}} \right] \exp \left[ -\frac{1}{2} \sum_{i=1}^n w_i (Y_i - \beta_0 - \beta_1 X_{i1} - \dots - \beta_{p-1} X_{i,p-1})^2 \right]$$

Minimize

# In matrix notation

$$W = \begin{bmatrix} w_1 & 0 & \dots & 0 \\ 0 & w_2 & \dots & 0 \\ \vdots & \vdots & & \vdots \\ 0 & 0 & \dots & w_n \end{bmatrix}$$

Normal Equation

$$(X^T W X) \hat{\beta}_w = X^T W Y$$

Estimators

$$\hat{\beta}_w = (X^T W X)^{-1} X^T W Y$$

Variance of the estimators

$$Var(\hat{\beta}_w) = (X^T W X)^{-1}$$

# Properties of the estimators

- Unbiased
- Consistent
- Minimum variance among unbiased linear estimators
- When weights are known  $\text{Var}(\hat{\beta}_w)$  is generally less than  $\text{Var}(\hat{\beta})$

# Error variances unknown

- We need to estimate the error variances.
- Residuals from an OLS gives valuable information about the error variances
- Two methods:
  - Estimation of variance function
  - Use of replicates or near replicates

# Estimation of variance

- Squared residual  $\hat{e}^2$  is an estimator of  $\sigma_i^2$
- Absolute residual  $|\hat{e}|$  is an estimator for  $\sigma_i$
- Idea
  - We can estimate the variance function describing the relation of  $\sigma_i^2$  to relevant predictor variables by first fitting the regression model using unweighted least squares and then regressing  $\hat{e}^2$  or  $|\hat{e}|$  against the appropriate predictor variables.
- $|\hat{e}|$  is preferred if outliers exist.

# General guidelines

1. A residual plot against  $X_l$  exhibits a megaphone shape.  $\Rightarrow$  Regress the absolute residuals against  $X_l$
2. A residual plot against  $\hat{Y}$  exhibits a megaphone shape.  $\Rightarrow$  Regress the absolute residuals against  $\hat{Y}$
3. A plot of the squared residuals against  $X_l$  exhibits an upward tendency.  $\Rightarrow$  Regress the squared residuals against  $X_l$
4. A plot of the squared residuals against  $X_l$  suggests that the variance increases rapidly with increases in  $X_l$  up to a point and then increases more slowly.  $\Rightarrow$  Regress the absolute residuals against  $X_l$  and  $X_l^2$ .

# What next?

After the variance function or the standard deviation function is estimated, the fitted values from this function are used to obtain the estimated weights:

$$w_i = \frac{1}{\hat{s}_i^2} \quad \text{where } \hat{s}_i \text{ is fitted value from standard deviation function}$$

$$w_i = \frac{1}{\hat{v}_i} \quad \text{where } \hat{v}_i \text{ is fitted value from variance function}$$

The parameters are then estimated as

$$\hat{\beta}_w = (X^T W X)^{-1} X^T W Y$$

# Use of Replicates or Near Replicates

- In designed experiments  $\sigma_i^2$  is estimated using replicate observations at each combination of levels of the predictor variables.
- In observation studies, near replicates may be used.
- For example, if the residual plot against  $X_l$  shows a megaphone appearance, cases with  $X_1$  values can be grouped together and the variance of the residuals in each group calculated.
  - The reciprocal of these variances are the weights.



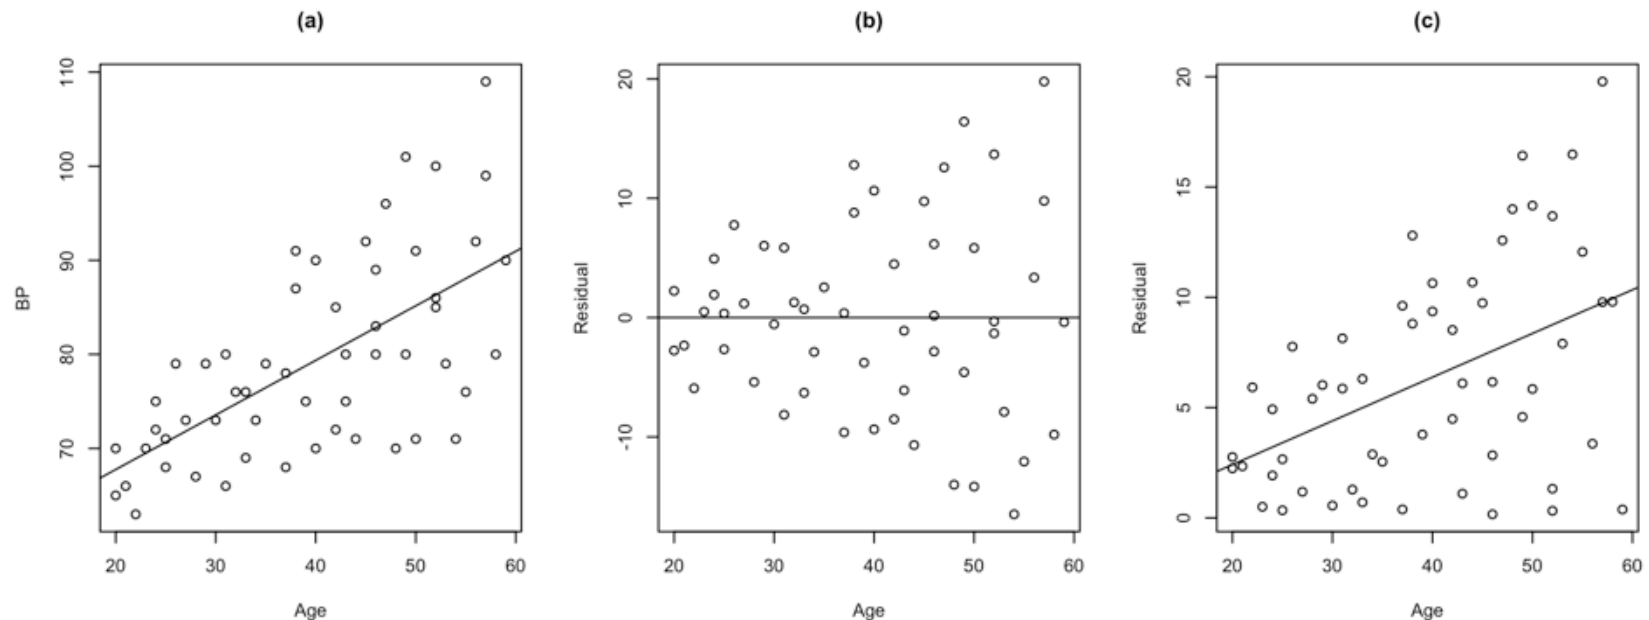
# Example (Strong Interaction ALR4 page 157)

- Response : scattering cross-section ( $y$ ), Predictor: square of total energy in the center of mass frame of reference ( $s$ )
- Designed experiment
- A very large number of particles was counted at each setting of  $s$
- The variance of  $y$  is thus known almost exactly

# Example (Strong Interaction ALR4 page 157)

```
##
## Call:
## lm(formula = y ~ x, data = alr4::physics, weights = 1/SD^2)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3230 -0.8842  0.0000  1.3900  2.3353
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   148.473      8.079   18.38 7.91e-08 ***
## x             530.835     47.550   11.16 3.71e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.657 on 8 degrees of freedom
## Multiple R-squared:  0.9397, Adjusted R-squared:  0.9321
## F-statistic: 124.6 on 1 and 8 DF, p-value: 3.71e-06
```

# Example (Blood pressure)



1.  $\Rightarrow$  linear relationship (unweighted)
2.  $\Rightarrow$  confirms the nonconstant error variance
3.  $\Rightarrow$  a linear relation between Age and standard error is reasonable

# Example (Blood pressure)

- Regress absolute residuals against Age

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.54948    2.18692 -0.7085 0.4817858
## blood$Age    0.19817    0.05309  3.7328 0.0004705
##
## n = 54, p = 2, Residual SE = 4.46057, R-Squared = 0.21
```

- Variance function

$$\hat{s} = -1.5494776 + 0.1981723Age$$

- Weights

```
w = 1/lmod_abs_res$fitted.values^2
head(w)

##           1           2           3           4           5           6
## 0.06920928 0.14655708 0.12661657 0.09725115 0.08625993 0.11048521

head(blood$Age)

## [1] 27 21 22 24 25 23
```

# Example (Blood pressure)

- OLS

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 56.156929   3.993674 14.0615 < 2.2e-16
## Age         0.580031    0.096951  5.9827 2.05e-07
##
## n = 54, p = 2, Residual SE = 8.14575, R-Squared = 0.41
```

- WLS

```
wls_mod = lm(BP ~ Age, weights = w, data = blood)
sumary(wls_mod)
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 55.565766   2.520918 22.042 < 2.2e-16
## Age         0.596342    0.079238  7.526 7.187e-10
##
## n = 54, p = 2, Residual SE = 1.21302, R-Squared = 0.52
```