

# Introduction to Python

Subrata Paul

6/3/2020

# Data types

```
l = [1,2,3,4]
t = (1,2,3,4)
s = {1,2,3,4}
d = {'a':'something', 'b':'something else'}
type(l)
```

```
## <class 'list'>
```

```
type(t)
```

```
## <class 'tuple'>
```

```
type(s)
```

```
## <class 'set'>
```

```
type(d)
```

```
## <class 'dict'>
```

# The range function

```
range(1,10)
```

```
## range(1, 10)
```

```
type(range(1,10))
```

```
## <class 'range'>
```

```
list(range(1,10))
```

```
## [1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
range(10)
```

```
## range(0, 10)
```

```
list(range(10))
```

```
## [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
list(range(1,10,2))
```

```
## [1, 3, 5, 7, 9]
```

# Libraries

```
import numpy as np  
np.linspace(2,3,5)
```

```
## array([2. , 2.25, 2.5 , 2.75, 3.  ])
```

```
np.arange(1,10)
```

```
## array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
np.sqrt(100)
```

```
## 10.0
```

```
x = np.array([1,2,3,4])  
np.sqrt(x)
```

```
## array([1.         , 1.41421356, 1.73205081, 2.         ])
```

```
np.sum(x)
```

```
## 10
```

# The **rep** function in R

```
[1,2] * 3
```

```
## [1, 2, 1, 2, 1, 2]
```

```
[1,2,3] + [2,3,4]
```

```
#[1,2,3] + 1 Does not work
```

```
## [1, 2, 3, 2, 3, 4]
```

```
np.array([1,2,3]) + 1
```

```
## array([2, 3, 4])
```

# Logical

```
x = np.random.normal(loc = 5, scale = 3, size = 10)
```

```
x
```

```
## array([-1.28484483, 10.33549025,  3.32330382,  5.1917064 ,  6.41168262,  
##          9.7346145 ,  5.26552406,  3.00029804,  6.90064899, 10.31778038])
```

```
x > 5
```

```
## array([False,  True, False,  True,  True,  True,  True, False,  True,  
##          True])
```

# Commonly-used Functions

```
x = np.random.normal(loc = 5, scale = 3, size = 10)
np.sum(x)
```

```
## 50.76907661221843
```

```
np.std(x)
```

```
## 3.8053914459548155
```

```
np.abs(x)
```

```
## array([ 7.52754874,  7.29607673,  4.95793231,  5.85435887,  1.02213247,
##         0.56686051,  9.02324184, 11.66286083,  3.67282345,  1.2295058 ])
```

```
np.var(x)
```

```
## 14.481004056946082
```

```
np.log(x)
```

```
## array([ 2.01856946,  1.98733677,  1.60098878,  1.76718649,          nan,
##        -0.56764203,  2.19980368,  2.4564095 ,  1.3009607 ,  0.2066123 ])
```

```
## /Users/spaul/Library/r-miniconda/envs/r-reticulate/bin/python:1: RuntimeWarning: invalid va
```

# Function

```
def five_point_summary(x):  
    mini = np.min(x)  
    maxi = np.max(x)  
    q = np.quantile(x, [.25, .50, .75])  
    return [mini] + list(q) + [maxi]  
x = np.random.normal(loc = 5, scale = 3, size = 10)  
five_point_summary(x)
```

```
## [-1.6936253268348889, 0.5257096814562654, 4.577801774829398, 6.725092755021153, 10.098419108]
```



# Functions related to statistical distribution

```
import scipy.stats as st
#loc = mean, scale = sd
st.norm.cdf(3, loc = 1, scale = 2) #pnorm
```

```
## 0.8413447460685429
```

```
st.norm.pdf(3, loc = 1, scale = 2) #dnorm
```

```
## 0.12098536225957168
```

```
st.norm.ppf(0.8413, loc = 1, scale = 2) # qnorm
```

```
## 2.9996301872294895
```

```
st.norm.rvs(loc =1, scale = 2, size = 10) # rnorm
```

```
## array([ 2.45493967,  4.10533824,  1.76931079,  0.78393315,  3.24838318,
##          0.8134111 ,  0.85946704, -0.55453708,  1.15296128,  0.24683208])
```

```
st.norm.rvs(10) # Standard normal
```

```
## 10.833291650203595
```

```
st.norm.rvs(size = 10)
```

```
## array([-0.92107289, -0.88847123, -0.83235645, -0.83131681,  1.29286413,
##          0.8134111 ,  0.85946704, -0.55453708,  1.15296128,  0.24683208])
```

# Data Frames

```
import pandas as pd
mydataframe = pd.DataFrame({'d':[1,2,3,4], 'e':['red','white','blue',np.NaN], 'f':[True, True,
mydataframe
```

```
##      d      e      f
## 0  1    red   True
## 1  2  white   True
## 2  3   blue  False
## 3  4    NaN   True
```

# Rename columns

```
mydataframe.columns = [ 'ID', 'Color', 'Passed' ]  
mydataframe
```

|      | ID | Color | Passed |
|------|----|-------|--------|
| ## 0 | 1  | red   | True   |
| ## 1 | 2  | white | True   |
| ## 2 | 3  | blue  | False  |
| ## 3 | 4  | NaN   | True   |

# Access columns

```
mydataframe.ID
```

```
## 0    1
## 1    2
## 2    3
## 3    4
## Name: ID, dtype: int64
```

```
mydataframe[['ID', 'Color']]
```

```
##      ID  Color
## 0     1    red
## 1     2  white
## 2     3   blue
## 3     4    NaN
```

# Indexing and Slicing

```
mydataframe.loc[1:2,['ID','Passed']]
```

```
##      ID  Passed
## 1     2     True
## 2     3    False
```

```
mydataframe.loc[:, 'ID']
```

```
## 0     1
## 1     2
## 2     3
## 3     4
## Name: ID, dtype: int64
```

# Indexing and Slicing

```
mydataframe.iloc[1:2,0]
```

```
## 1      2  
## Name: ID, dtype: int64
```

```
mydataframe.iloc[:,[0,1]]
```

```
##      ID  Color  
## 0     1    red  
## 1     2  white  
## 2     3   blue  
## 3     4    NaN
```

# Importing Data

```
help(pd.read_csv)
```

```
## Help on function read_csv in module pandas.io.parsers:
##
## read_csv(filepath_or_buffer, sep=',', delimiter=None, header='infer', names=None, index_col=
##      Read a comma-separated values (csv) file into DataFrame.
##
##      Also supports optionally iterating or breaking of the file
##      into chunks.
##
##      Additional help can be found in the online docs for
##      `IO Tools <http://pandas.pydata.org/pandas-docs/stable/io.html>`_.
##
## Parameters
## -----
## filepath_or_buffer : str, path object, or file-like object
##      Any valid string path is acceptable. The string could be a URL. Valid
##      URL schemes include http, ftp, s3, and file. For file URLs, a host is
##      expected. A local file could be: file://localhost/path/to/table.csv.
##
##      If you want to pass in a path object, pandas accepts either
##      ``pathlib.Path`` or ``py._path.local.LocalPath``.
```

# Sample data analysis

```
pima = pd.read_csv('./data/pima.txt', header = 0, sep = '\t')  
pima.head()
```

| ##   | pregnant | glucose | diastolic | triceps | insulin | bmi  | diabetes | age | test |
|------|----------|---------|-----------|---------|---------|------|----------|-----|------|
| ## 0 | 6        | 148     | 72        | 35      | 0       | 33.6 | 0.627    | 50  | 1    |
| ## 1 | 1        | 85      | 66        | 29      | 0       | 26.6 | 0.351    | 31  | 0    |
| ## 2 | 8        | 183     | 64        | 0       | 0       | 23.3 | 0.672    | 32  | 1    |
| ## 3 | 1        | 89      | 66        | 23      | 94      | 28.1 | 0.167    | 21  | 0    |
| ## 4 | 0        | 137     | 40        | 35      | 168     | 43.1 | 2.288    | 33  | 1    |



# Assign NULL for missing data

```
pima.diastolic.sort_values().head()
```

```
## 347    0
## 494    0
## 222    0
## 81     0
## 78     0
## Name: diastolic, dtype: int64
```

```
pima.diastolic[pima.diastolic==0] = np.NaN
```

```
## /Users/spaul/Library/r-miniconda/envs/r-reticulate/bin/python:1: SettingWithCopyWarning:
## A value is trying to be set on a copy of a slice from a DataFrame
##
## See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html
```

```
pima.diastolic.sort_values(na_position = 'first').head()
```

```
## 7      NaN
## 15     NaN
## 49     NaN
## 60     NaN
## 78     NaN
## Name: diastolic, dtype: float64
```

# Assign NaN to missing data

```
pima.glucose[pima.glucose==0] = np.NaN  
pima.triceps[pima.triceps==0] = np.NaN  
pima.insulin[pima.insulin==0] = np.NaN  
pima.bmi[pima.bmi==0] = np.NaN
```

# Assign NaN to missing data

```
pima = pd.read_csv('./data/pima.txt', header = 0, sep = '\t')
pima.replace(0,np.nan, inplace = True)
pima.head()
```

| ##   | pregnant | glucose | diastolic | triceps | insulin | bmi  | diabetes | age | test |
|------|----------|---------|-----------|---------|---------|------|----------|-----|------|
| ## 0 | 6.0      | 148.0   | 72.0      | 35.0    | NaN     | 33.6 | 0.627    | 50  | 1.0  |
| ## 1 | 1.0      | 85.0    | 66.0      | 29.0    | NaN     | 26.6 | 0.351    | 31  | NaN  |
| ## 2 | 8.0      | 183.0   | 64.0      | NaN     | NaN     | 23.3 | 0.672    | 32  | 1.0  |
| ## 3 | 1.0      | 89.0    | 66.0      | 23.0    | 94.0    | 28.1 | 0.167    | 21  | NaN  |
| ## 4 | NaN      | 137.0   | 40.0      | 35.0    | 168.0   | 43.1 | 2.288    | 33  | 1.0  |

# Plotting Histogram

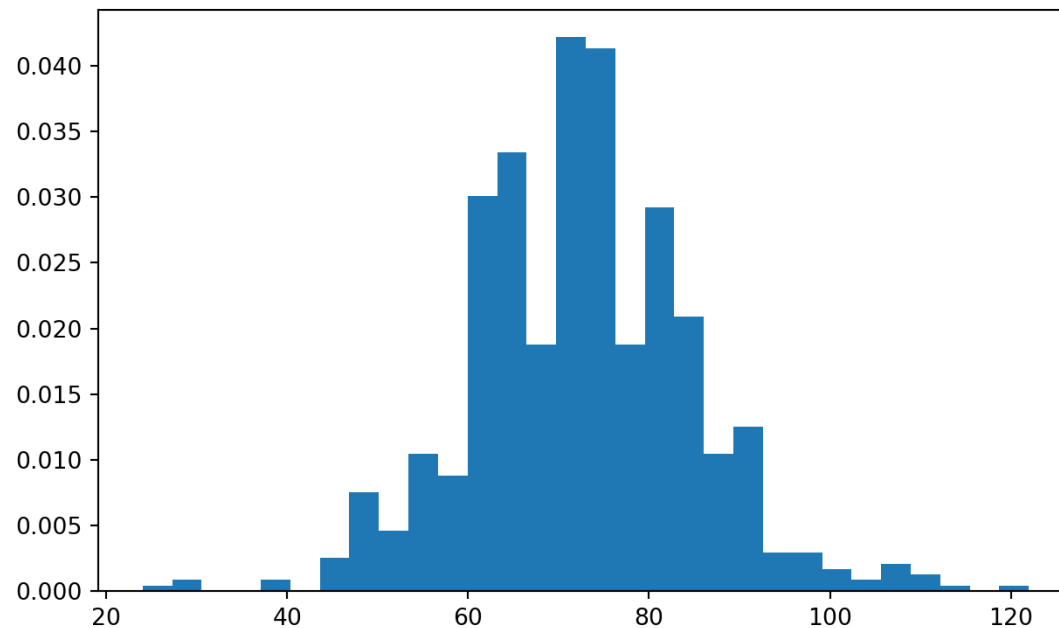
```
import matplotlib.pyplot as plt  
_ = plt.hist(pima.diastolic)
```

```
## /Users/spaul/Library/r-miniconda/envs/r-reticulate/lib/python3.6/site-packages/numpy/lib/his  
## keep = (tmp_a >= first_edge)  
## /Users/spaul/Library/r-miniconda/envs/r-reticulate/lib/python3.6/site-packages/numpy/lib/his  
## keep &= (tmp_a <= last_edge)
```

```
plt.show()
```

# Plotting Histogram

```
import matplotlib.pyplot as plt
_ = plt.hist(pima.diastolic[~np.isnan(pima.diastolic)], bins = 30, density = True)
plt.show()
```



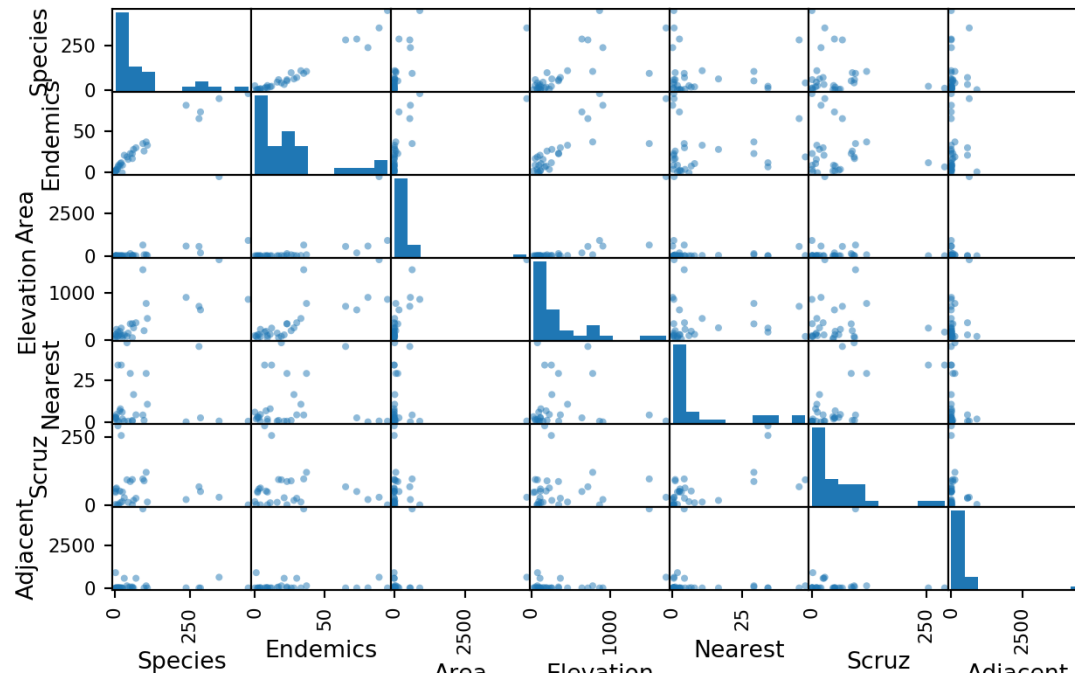
# Galapagos Data

```
gala = pd.read_csv('./data/gala.txt', header = 0, sep = '\t')
gala.head()
```

| ##           | Species | Endemics | Area  | Elevation | Nearest | Scruz | Adjacent |
|--------------|---------|----------|-------|-----------|---------|-------|----------|
| ## Baltra    | 58      | 23       | 25.09 | 346       | 0.6     | 0.6   | 1.84     |
| ## Bartolome | 31      | 21       | 1.24  | 109       | 0.6     | 26.3  | 572.33   |
| ## Caldwell  | 3       | 3        | 0.21  | 114       | 2.8     | 58.7  | 0.78     |
| ## Champion  | 25      | 9        | 0.10  | 46        | 1.9     | 47.4  | 0.18     |
| ## Coamano   | 2       | 1        | 0.05  | 77        | 1.9     | 1.9   | 903.82   |

# Plot data

```
_ = pd.plotting.scatter_matrix(gala)  
plt.show()
```



# Fit multiple regression model

```
import statsmodels.api as sm
Y = gala.Species
X = gala.iloc[:,2:]
X['const'] = np.ones(X.shape[0])
lmod = sm.OLS(Y, X)
res = lmod.fit()
print(res.summary())
```

```
##                                OLS Regression Results
## =====
## Dep. Variable:                Species    R-squared:                0.766
## Model:                        OLS        Adj. R-squared:         0.717
## Method:                      Least Squares    F-statistic:           15.70
## Date:                        Fri, 17 Jul 2020    Prob (F-statistic):    6.84e-07
## Time:                        17:32:54        Log-Likelihood:        -162.54
## No. Observations:            30            AIC:                   337.1
## Df Residuals:                24            BIC:                   345.5
## Df Model:                    5
## Covariance Type:             nonrobust
## =====
##                                coef    std err          t      P>|t|    [0.025    0.975]
## -----
```



# Fit multiple regression model

```
import statsmodels.api as sm
Y = gala.Species
X = gala.iloc[:,2:]
X['const'] = np.ones(X.shape[0])
lmod = sm.OLS(Y, X)
res = lmod.fit()
print(res.summary())
```

```
##                                OLS Regression Results
## =====
## Dep. Variable:                Species    R-squared:                0.766
## Model:                        OLS        Adj. R-squared:           0.717
## Method:                      Least Squares    F-statistic:             15.70
## Date:                        Fri, 17 Jul 2020    Prob (F-statistic):       6.84e-07
## Time:                        17:32:54        Log-Likelihood:           -162.54
## No. Observations:            30            AIC:                     337.1
## Df Residuals:                24            BIC:                     345.5
## Df Model:                    5
## Covariance Type:             nonrobust
## =====
##                                coef    std err          t      P>|t|      [0.025      0.975]
## -----
```

# Parameters

```
res.params
```

```
## Area      -0.023938
## Elevation  0.319465
## Nearest    0.009144
## Scrutz    -0.240524
## Adjacent  -0.074805
## const      7.068221
## dtype: float64
```

# Standard Errors

```
res.bse
```

```
## Area          0.022422
## Elevation     0.053663
## Nearest       1.054136
## Scruz         0.215402
## Adjacent      0.017700
## const        19.154198
## dtype: float64
```

# Fitted Values

res.fittedvalues

|                 |            |
|-----------------|------------|
| ## Baltra       | 116.725946 |
| ## Bartolome    | -7.273154  |
| ## Caldwell     | 29.330659  |
| ## Champion     | 10.364266  |
| ## Coamano      | -36.383916 |
| ## Daphne.Major | 43.087705  |
| ## Daphne.Minor | 33.919668  |
| ## Darwin       | -9.018992  |
| ## Eden         | 28.314202  |
| ## Enderby      | 30.785943  |
| ## Espanola     | 47.656487  |
| ## Fernandina   | 96.989598  |
| ## Gardner1     | -4.033276  |
| ## Gardner2     | 64.633796  |
| ## Genovesa     | -0.497176  |
| ## Isabela      | 386.403558 |
| ## Marchena     | 88.694540  |
| ## Onslow       | 4.037233   |
| ## Pinta        | 215.679486 |
| ## Pinzon       | 150.475375 |